

Sample Engineering Background

Quill Healey

June 10, 2024

Contents

1	Estimation, Fitting, Detection, and Classification	3
1.1	Markov Chain Transition Probability Estimation	3
1.2	Generalized Additive Model	7
1.3	Bounding Object Position	10
1.4	Robust Logistic Regression	14
2	Design and Control	17
2.1	Minimum Fuel Control	17
2.2	Path Planning with Contingencies	22
2.3	Output Tracking	25
2.4	UAV Design	29
2.5	Optimal Spacecraft Landing	36

Document Overview

This collection of problems is designed to demonstrate the engineering and mathematics background that I developed during my time as an undergraduate student at Georgia Tech (and continue to develop now on my own). While every example is a “toy problem” in the sense that one, they all have associated solutions (thus differing from research), and two, this solution can be obtained within a few pages of work (thus differing from a larger-scope engineering project), I chose to devote time to this pursuit as **I believe this collection of problems demonstrates the variety of ways I could be of use to Anduril as a member on their Mission Autonomy or Perception teams.**

Each problem contains three parts. The first is a high-level overview of the problem; it is meant to be understood by people from non-technical backgrounds. The length of this description varies according to the complexity of the problem it is trying to summarize. The description to Problem 1.1 is the longest and most technically involved, so if this is discouraging, move to another problem; each problem stands on its own. In fact, I recommend the reader use the hyperlinked table of contents to explore the problem which seems most interesting to them.

The second and third part of each problem are mathematically involved. They are, respectively, the prompt of a graduate level problem (taken from one of [BV04], [Boy-8], or [BV24]) and my response/solution to that problem. However, while these two parts are less approachable to non-technical readers, my response to almost every one contains some figure or table, which when connected with the first part of the problem still yields a clear (and in my opinion, compelling) depiction of my skillset.

Finally, some housekeeping.

- [Here](#) is the repository containing this document, the larger note set that this document is a part of, and all of the code used to solve problems in this document. The code can be found in the “examples” subdirectory.
- There is not much exotic notation in this document. I refer a reader to [BV04] for a summary of all notation that I use (I have adopted the style from that book). The one bit of notation which I’ve found confuses people are the generalized inequalities \succeq and \preceq . In this document (as it currently exists) these are just used to denote *componentwise* inequality between vectors.

1 Estimation, Fitting, Detection, and Classification

1.1 Markov Chain Transition Probability Estimation

Problem Overview

Markov chains are an important, probabilistic tool that can be found across many disciplines, either used to model some system directly (like in the following example) or as part of a broader solution approach (*e.g.* Markov Chain Monte Carlo or Markov Decision Processes). We'll forego the mathematics behind Markov chains and consider the canonical weather example to build intuition for this problem.

Suppose on any given day it can be sunny, cloudy, or rainy. These are known as the *states* of the Markov chain. The important assumption we make is that the **weather tomorrow only depends on the weather today**. However, given the weather today, we do not know for sure what the weather tomorrow will be, *i.e.* we transition randomly between states. Figure 1 is an example of a Markov chain that models the weather. According to this model, if it is sunny today, it will be rainy tomorrow with probability 0.1, cloudy tomorrow with probability 0.3, or sunny again tomorrow with probability 0.6.




Figure 1: Markov chain weather model

The problem in this section is concerned with *estimating* the transition probabilities between states having *observed* a sequence of states. In other words, it is a **statistical question**. Perhaps we wish to create a new weather app because we believe that our Markov model will do a better job *predicting* the weather than the default app on our phones (a bold assumption). However, while we know what the states in our system are, we do not know what the transition probabilities are; *i.e.*, our knowledge of the system is represented by Figure 2. We could of course *guess* what each probability is, but the better approach is to *collect data* and then use this data to *estimate* the probabilities (fun fact: statistics used to be called *inverse probability*).

For instance, perhaps we kept track of the weather over the past week and observed that on Monday it rained, on Tuesday it was cloudy, on Wednesday it was sunny, sunny again on Thursday, cloudy on Friday, and then it rained on both Saturday and Sunday. The




Figure 2: Markov chain weather model with unknown transition probabilities

methodology introduced in the problem below would be one way of estimating the probability links in figure 2's state diagram.

Notes:

- This problem uses what Gilbert Strang would refer to as a *Markov Matrix*. In more traditional, statistical texts, the transpose of this matrix is worked with, and it is called a *stochastic matrix*.
- There are many other methodologies for approaching this inference problem, almost all of which are more accurate the majority of the time. However, this frequentist maximum likelihood approach is often still used as part of these other, “better” methodologies.

Problem

[BV04] **Exercise 7.5.** *Markov chain estimation.* Consider a Markov chain with n states and a transition probability matrix $P \in \mathbf{R}^{n \times n}$ defined as

$$P_{ij} = \mathbf{Prob}(y_{t+1} = i \mid y_t = j).$$

The transition probabilities must satisfy $P_{ij} \geq 0$ and $\sum_{i=1}^n P_{ij} = 1, j = 1, \dots, n$. We consider the problem of estimating the transition probabilities, given an observed sample sequence $y_1 = k_1, y_2 = k_2, \dots, y_N = k_n$.

(a) Show that if there are no other prior constraints on P_{ij} , then the ML estimates are the empirical transition frequencies: \hat{P}_{ij} is the ratio of the number of times the state transitioned from j into i , divided by the number of times it was j , in the observed sample.

(b) Suppose that an equilibrium distribution p of the Markov chain is known, i.e., a vector $q \in \mathbf{R}_+^n$ satisfying $\mathbf{1}^T q = 1$ and $Pq = q$. Show that the problem of computing the ML estimate of P , given the observed sequence and knowledge of q , can be expressed as a convex optimization problem.

My Response

(a) Firstly, we'll let

$$p_P(k) = \mathbf{Prob}_{y \sim P}(y_1 = k_1, \dots, y_N = k_n)$$

denote either the *joint probability density function*, or, the *likelihood function*. Which it is depends on context. Since we are in a statistical setting, it is the likelihood function. Furthermore, we seek to find an estimate of the transition probabilities between the states of our Markov chain, *i.e.*, we wish to find $\hat{P} \in \mathbf{R}^{n \times n}$. As specified above, for this problem we will find this estimate using maximum likelihood estimation, corresponding to

$$\hat{P} = \underset{P}{\operatorname{argmax}} p_P(k).$$

We first expand the likelihood function using *the chain rule* (abbreviate $x_i = k_i$ to just x_i)

$$p_P(k) = \mathbf{Prob}(x_N | x_{N-1}, \dots, x_1) \mathbf{Prob}(x_{N-1} | x_{N-2}, \dots, x_1) \cdots \mathbf{Prob}(x_2 | x_1) \mathbf{Prob}(x_1)$$

and then use our *Markov* assumption to simplify this expression to

$$p_P(k) = P_{k_n k_{n-1}} P_{k_{n-1} k_{n-2}} \cdots P_{k_2 k_1},$$

which is simply a product of probabilities, *i.e.*, $P_{ij} \in \mathbf{R}$ (and we assume that we start in state k_1). Letting c_{ij} be the *number of times the state transitioned from j to i* , the likelihood function becomes

$$p_P(k) = \prod_{i,j=1}^n P_{ij}^{c_{ij}}.$$

For tractability purposes we now make use of the *log-likelihood function*, which is simply

$$l(P) = \log \left(\prod_{i,j=1}^n P_{ij}^{c_{ij}} \right) = \sum_{i,j=1}^n c_{ij} \log P_{ij}.$$

To find the transition probability matrix P which maximizes the likelihood of our observed data having occurred, we thus want to solve

$$\begin{array}{ll} \text{maximize} & \sum_{i,j=1}^n c_{ij} \log P_{ij} \\ \text{subject to} & \mathbf{1}^T P = \mathbf{1}^T. \end{array}$$

(It is critical to formulate this *constrained problem*. Forgetting to include the constraint is a common mistake which leads to an underdetermined problem.)

To solve this equality-constrained optimization problem we simply form the Lagrangian

$$L(P, \nu) = \sum_{i,j=1}^n c_{ij} \log P_{ij} + \sum_{j=1}^n \left(\nu_j \sum_{i=1}^n P_{ij} - 1 \right),$$

find the partial derivative of L with respect to some probability, P_{ij} ,

$$\frac{\partial L}{\partial P_{ij}}(P, \nu) = \frac{c_{ij}}{P_{ij}} + v_j$$

and set $\frac{\partial L}{\partial P_{ij}}(P, \nu)$ equal to zero:

$$\frac{c_{ij}}{P_{ij}} = -v_j \iff P_{ij} = \frac{-c_{ij}}{v_j}.$$

Then, plugging this expression for P_{ij} into the constraint

$$\mathbf{1}^T P = \mathbf{1}^T \iff \sum_{i=1}^n P_{ij} = 1, \quad j = 1, \dots, n,$$

we find that

$$\sum_{i=1}^n \frac{-c_{ij}}{v_j} = 1, \quad j = 1, \dots, n \iff v_j = \sum_{i=1}^n -c_{ij}, \quad j = 1, \dots, n.$$

Therefore, our final expression for any transition probability is

$$\hat{P}_{ij} = \frac{-c_{ij}}{v_j} = \frac{-c_{ij}}{\sum_{i=1}^n -c_{ij}} = \frac{c_{ij}}{\sum_{i=1}^n c_{ij}},$$

which is of course the empirical transition probability, as desired.

Note that since this is a convex optimization problem (concave objective function constrained to a convex set) we don't need to check boundary conditions.

(b) To incorporate additional information about the transition probabilities that we are trying to estimate, we formulate

$$\begin{aligned} & \text{maximize} && \sum_{i,j=1}^n c_{ij} \log P_{ij} \\ & \text{subject to} && \mathbf{1}^T P = \mathbf{1} \\ & && Pq = q, \end{aligned}$$

which is still a convex problem. If we wanted to, we could augment the above Lagrangian with these additional constraints and find an analytical expression for this problem, but since it is not required, we can be satisfied knowing that CVXPY could solve this problem with even non-trivial scale data in under a second.

1.2 Generalized Additive Model

Problem Overview

This problem is straightforward (albeit the model is less so): we wish to *guess* the value of a function which happens to be some combination of other functions and additional *randomness*. This is the setup of modern machine learning. There are many other ML models that will give a better estimate of the underlying function, but a benefit to this model is that the problem being solved to create it is *convex*, which means that we could embed the corresponding software into an online system that observes and reacts to data that the function produces without worrying about the software failing.

Problem

[BV24] **Exercise 6.17.** *Fitting a generalized additive regression model. A generalized additive model has the form*

$$f(x) = \alpha + \sum_{j=1}^n f_j(x_j)$$

for $x \in \mathbf{R}^n$, where $\alpha \in \mathbf{R}$ is the offset, and $f_j : \mathbf{R} \rightarrow \mathbf{R}$, with $f_j(0) = 0$. The functions f_j are called the regressor functions. When each f_j is linear, i.e., has the form $w_j x_j$, the generalized additive model is the same as the standard (linear) regression model. Roughly speaking, a generalized additive model takes into account nonlinearities in each regressor x_j , but not nonlinear interactions among the regressors. To visualize a generalized additive model, it is common to plot each regressor function (when n is not too large). We will restrict the functions f_j to be piecewise-affine, with given knot points $p_1 < \dots < p_K$. This means that f_j is affine on the intervals $(-\infty, p_1], [p_1, p_2], \dots, [p_{K-1}, p_K], [p_K, \infty)$, and continuous at p_1, \dots, p_K . Let C denote the total (absolute value of) change in slope across all regressor functions and all knot points. The value C is a measure of nonlinearity of the regressor functions; when $C = 0$, the generalized additive model reduces to a linear regression model. Now suppose we observe samples or data $(x^{(1)}, y^{(1)}), \dots, (x^{(N)}, y^{(N)}) \in \mathbf{R}^n \times \mathbf{R}$, and wish to fit a generalized additive model to the data. We choose the offset and the regressor functions to minimize

$$\frac{1}{N} \sum_{i=1}^N (y^{(i)} - f(x^{(i)}))^2 + \lambda C.$$

where $\lambda > 0$ is a regularization parameter. (The first term is the mean-square error.)

(a) Explain how to solve this problem using convex optimization.

(b) Carry out the method of part (a) using the data in the file `gen_add_reg_data.m`. This file contains the data, given as an $N \times n$ matrix X (whose rows are $(x^{(i)})^T$), a column vector y (which give $y^{(i)}$), a vector p that gives the knot points, and the scalar λ . Give the mean-square error achieved by your generalized additive regression model. Compare the estimated and true regressor functions in a 3×3 array of plots (using the plotting code in the data file as a template), over the range $-10 \leq x_i \leq 10$. The true regressor functions (to be used only for plotting, of course) are given in the cell array f .

My Response

(a) The critical step to using convex optimization to solve this problem is recognizing that we wish to find the value of $\theta \in \mathbf{R}^{n(K+1)}$, where these function parameters are found in the feature mappings according to

$$f_j(x) = \sum_{k=0}^K \theta_{nk+j} b_k(x),$$

where $b_k(x) = \max\{x - p_k, 0\} + \min\{p_k, 0\}$. Furthermore, letting

$$C = \sum_{j=1}^{n(K+1)} |\theta_j| = \|\theta\|_1,$$

the following unconstrained optimization problem

$$\text{minimize} \quad \frac{1}{N} \sum_{i=1}^N \left(y^{(i)} - \left(\alpha + \sum_{j=1}^n f_j(x_j^{(i)}) \right) \right)^2 + \lambda C,$$

is convex.

(b) The code in algorithm 1 carries out this method, with the key functionality coming from the augmented data matrix (numpy matrix Z). The fit generalized additive model has a MSE of 0.96962. A comparison of the estimated and true regressor functions can be found in figure 3.

Algorithm 1: Generalized additive regression model

```
1      Z = X.copy()
2      for i in range(K):
3          Z = np.hstack((Z, np.maximum(X-p[i], 0) + np.minimum(p[i],
4              0)))
5
6      theta = cp.Variable(n*(K+1))
7      alpha = cp.Variable(1)
8
9      C = cp.norm(theta, 1)
10     f0 = 1/N*cp.sum_squares(y - alpha - Z@theta)
11
12     prob = cp.Problem(cp.Minimize(f0 + lambda_value*C), [])
13     prob.solve()
```




Figure 3: True versus Estimated Regressor Functions.

1.3 Bounding Object Position

Problem Overview

The objective of this problem is straightforward: we have some number of cameras observing a physical space. For illustrative purposes, perhaps these cameras are part of a *targeting system*. Furthermore, suppose that there's an object in this physical space that we wish to terminate. Therefore, we need to identify the physical coordinates that this object occupies. The difficulty is that while the object is in three dimensional space, the image plane is two dimensional. The following problem dives into the mathematical relationship between image space and physical space to create the tools necessary to use the images that the cameras take to find the smallest possible box in physical space which this object occupies.

Problem

[BV24]**Exercise 8.8.** *Bounding object position from multiple camera views.* $x \in \mathbf{R}^3$, and viewed by a set of m cameras. Our goal is to find a box in \mathbf{R}^3 ,

$$\mathcal{B} = \{z \in \mathbf{R}^3 \mid l \preceq z \preceq u\},$$

for which we can guarantee $x \in \mathcal{B}$. We want the smallest possible such bounding box. (Although it doesn't matter, we can use volume to judge 'smallest' among boxes.)

Now we describe the cameras. The object at location $x \in \mathbf{R}^3$ creates an image on the image plane of camera i at location

$$v_i = \frac{1}{c_i^T x + d_i} (A_i x + b_i) \in \mathbf{R}^2.$$

The matrices $A_i \in \mathbf{R}^{2 \times 3}$, vectors $b_i \in \mathbf{R}^2$ and $c_i \in \mathbf{R}^3$, and real numbers $d_i \in \mathbf{R}$ are known, and depend on the camera positions and orientations. We assume that $c_i^T x + d_i > 0$. The 3×4 matrix

$$P_i = \begin{bmatrix} A_i & b_i \\ c_i^T & d_i \end{bmatrix}$$

is called the camera matrix (for camera i). It is often (but not always) the case that the first 3 columns of P_i (i.e., A_i stacked above c_i^T) form an orthogonal matrix, in which case the camera is called orthographic. We do not have direct access to the image point v_i ; we only know the (square) pixel that it lies in. In other words, the camera gives us a measurement \hat{v}_i (the center of the pixel that the image point lies in); we are guaranteed that

$$\|v_i - \hat{v}_i\|_\infty \leq \rho_i/2,$$

where ρ_i is the pixel width (and height) of camera i . (We know nothing else about v_i ; it could be any point in this pixel.) Given the data $A_i, b_i, c_i, d_i, \hat{v}_i, \rho_i$, we are to find the smallest box \mathcal{B} (i.e., find the vectors l and u) that is guaranteed to contain x . In other words, find the smallest box in \mathbf{R}^3 that contains all points consistent with the observations from the camera.

- (a) Explain how to solve this using convex or quasiconvex optimization. You must explain any transformations you use, any new variables you introduce, etc. If the convexity or quasiconvexity of any function in your formulation isn't obvious, be sure justify it.
- (b) Solve the specific problem instance given in the file `camera_data.m`. Be sure that your final numerical answer (i.e., l and u) stands out.

My Response

Let's first formulate the constraints ensuring that the box that we propose x lies in is consistent with the camera measurements. Mathematically, this is expressed as

$$\|v_i - \hat{v}_i\|_\infty \leq \rho_i/2, \quad i = 1, \dots, m,$$

which can then be expanded according to v_i 's definition

$$\left\| \frac{1}{c_i^T x + d_i} (A_i x + b_i) - \hat{v}_i \right\|_\infty \leq \rho_i/2, \quad i = 1, \dots, m.$$

The ℓ_∞ -norm is requiring that the absolute value of both terms in the \mathbf{R}^2 vector

$$\frac{1}{c_i^T x + d_i} (A_i x + b_i) - \hat{v}_i$$

be less than or equal to $\rho_i/2$. Furthermore, we can re-express these constraints as

$$-(\rho_i/2)\mathbf{1} \preceq \frac{1}{c_i^T x + d_i} (A_i x + b_i) - \hat{v}_i \preceq (\rho_i/2)\mathbf{1}, \quad i = 1, \dots, m$$

and then as

$$(\hat{v}_i - (\rho_i/2)\mathbf{1})(c_i^T x + d_i) \preceq A_i x + b_i \preceq (\hat{v}_i + (\rho_i/2)\mathbf{1})(c_i^T x + d_i), \quad i = 1, \dots, m,$$

where we can multiply through by $c_i^T x + d_i$ as it is assumed positive. This is simply a set of $2m$ linear inequalities on the position x .

So, by using these $2m$ linear inequalities as constraints in an optimization problem we can ensure that a proposed position is validated by our pictures. To find the minimum box which x can lie in is now simple: Consider a single coordinate of physical space (for instance x_1 we can imagine as N/S, x_2 as W/E, and x_3 as altitude). To find the maximum northward point along the N/S line which is consistent with our measurements we simply select an x_1 which satisfies our $2m$ inequalities, and then maximize this scalar value until right before it violates a single inequality. That is informal speak for saying, solve the following optimization problem

$$\begin{aligned} & \text{minimize} && x_1 \\ & \text{subject to} && (\hat{v}_i - (\rho_i/2)\mathbf{1})(c_i^T x + d_i) \preceq A_i x + b_i, \quad i = 1, \dots, m \\ & && A_i x + b_i \preceq (\hat{v}_i + (\rho_i/2)\mathbf{1})(c_i^T x + d_i), \quad i = 1, \dots, m. \end{aligned}$$

The remaining five coordinates are found in the same manner.

(b) See algorithms 2, 3, 4 for an instantiation of this problem. The vectors defining the minimum bounding box are

$$l = (-0.9956, 0.2753, -0.6790), \quad u = (-0.8245, 0.3783, -0.5735).$$

Algorithm 2: Bounding Box Data

```
1      ### DATA ###
2      P1 = np.array([
3          [-0.7430, -0.2438, -0.6233, 20.0000],
4          [0.1724, -0.9696, 0.1737, -10.0000],
5          [-0.6467, 0.0216, 0.7625, 2.0000]
6      ])
7
8      P2 = np.array([
9          [-0.9247, 0.2152, -0.3140, -10.0000],
10         [0.3105, 0.9037, -0.2948, 10.0000],
11         [0.2204, -0.3701, -0.9025, 3.0000]
12     ])
13
14     P3 = np.array([
15         [-0.6026, -0.3108, -0.7350, 4.0000],
16         [-0.5440, -0.5139, 0.6633, 8.0000],
17         [-0.5839, 0.7995, 0.1406, 1.5000]
18     ])
19
20     P4 = np.array([
21         [-0.6700, 0.0028, -0.7424, -10.0000],
22         [-0.6916, 0.3612, 0.6255, 4.0000],
23         [0.2699, 0.9325, -0.2400, 4.0000]
24     ])
25
26     rho = 0.1 * np.ones((4, 1))
27
28     vhat = np.array([
29         [9.9100, -2.7400, 2.2200, -2.1200],
30         [-5.0000, 3.1400, 3.6000, 1.0300]
31     ])
32
```

Algorithm 3: Bounding Box Objects

```
1      # Create A_i, b_i, c_i, and d_i objects
2      m=4
3      A1 = P1[0:2, 0:3]
4      A2 = P2[0:2, 0:3]
5      A3 = P3[0:2, 0:3]
6      A4 = P4[0:2, 0:3]
7      A_i = [A1, A2, A3, A4]
8
9      b1 = P1[0:2, 3].reshape((2, 1))
10     b2 = P2[0:2, 3].reshape((2, 1))
11     b3 = P3[0:2, 3].reshape((2, 1))
12     b4 = P4[0:2, 3].reshape((2, 1))
13     b_i = [b1, b2, b3, b4]
14
15     c1_T = P1[2, 0:3].reshape((1, 3))
16     c2_T = P2[2, 0:3].reshape((1, 3))
17     c3_T = P3[2, 0:3].reshape((1, 3))
18     c4_T = P4[2, 0:3].reshape((1, 3))
19     cT_i = [c1_T, c2_T, c3_T, c4_T]
20
21     d1 = P1[2, 3]
22     d2 = P2[2, 3]
23     d3 = P3[2, 3]
24     d4 = P4[2, 3]
25     d_i = [d1, d2, d3, d4]
26
```

Algorithm 4: Bounding Box Solution

```
1      # solve 6 optimization problems
2      x = cp.Variable((3, 1))
3
4      constra = [A_i[i]@x + b_i[i] <= ( vhat[:, i] + rho[i][0]/2 ).
5      reshape((2, 1)) * (cT_i[i]@x + d_i[i] ) for i in range(m)]
6      constra += [A_i[i]@x + b_i[i] >= ( vhat[:, i] - rho[i][0]/2 ).
7      reshape((2, 1)) * (cT_i[i]@x + d_i[i] ) for i in range(m)]
8
9      for i in range(3):
10         obj1 = cp.Minimize(x[i, 0])
11         obj2 = cp.Maximize(x[i, 0])
12         prob1 = cp.Problem(obj1, constra)
13         prob2 = cp.Problem(obj2, constra)
14         prob1.solve()
15         prob2.solve()
16         print(f"l_{i+1} = {prob1.value}")
17         print(f"u_{i+1} = {prob2.value}")
```

1.4 Robust Logistic Regression

Problem Overview

Logistic regression is a common approach to solving classification problems. A good example of such a problem is in a medical setting. For instance, suppose we are a doctor and it is our job to determine if a patient has a disease. If they have the disease, it is advantageous to take medical action. If they do not have the disease, performing this medical action can be costly to the patient, both financially and to their overall health. Furthermore, it is critical that we can properly *classify* whether they have this disease. Logistic regression can be used to produce a probability estimate that the patient has the disease (note that in the problem below the logistic models just produce binary output; that is, in this example, whether the patient has the disease or not. It would take no work to turn this into a probability).

Now, here's a twist on the problem. While some of the “features” that we might use to classify the patient's condition might be known exactly (*e.g.*, their height), perhaps our classification model also depends on some sensor readings of the patient's internal organs. Our sensors can only be so accurate, so it is very likely that the readings obtained deviate from the true value we are trying to measure by some amount (mathematically, we'll call this small deviation ϵ). This *robust logistic regression* problem addresses this uncertainty in feature data.

Problem

[BV24] **Exercise 6.29.** *Robust Logistic Regression.* We are given a dataset $x_i \in \mathbf{R}^d$, $y_i \in \{-1, 1\}$, $i = 1, \dots, n$. We seek a prediction model $y \approx \hat{y} = \text{sign}(\theta^T x)$, where $\theta \in \mathbf{R}^d$ is the model parameter. In logistic regression, θ is chosen as the minimizer of the logistic loss

$$\ell(\theta) = \sum_{i=1}^n \log(1 + \exp(-y_i \theta^T x_i))$$

which is a convex function of θ . (We will assume that a minimizer exists.)

We will take into account the idea that the feature vectors x_i are not known precisely. Specifically we imagine that each entry of each feature vector can vary by $\pm\epsilon$, where $\epsilon > 0$ is a given uncertainty level. We define the worst-case logistic loss as

$$\ell^{\text{wc}}(\theta) = \sum_{i=1}^n \sup_{\|\delta_i\|_\infty \leq \epsilon} \log(1 + \exp(-y_i \theta^T (x_i + \delta_i))).$$

In words: we perturb each feature vector's entries by up to ϵ in such a way as to make the logistic loss as large as possible. Each term is convex, since it is the supremum of a family of convex functions of θ , and so $\ell^{\text{wc}}(\theta)$ is a convex function of θ .

In *robust logistic regression*, we choose θ to minimize $\ell^{\text{wc}}(\theta)$. (Here too we assume a minimizer exists.)

(a) Explain how to carry out robust logistic regression by solving a single convex optimization problem in disciplined convex programming (DCP) form. Justify any change of variables or

introduction of new variables. Explain why solving the problem you propose also solves the robust logistic regression problem.

Hint: $\log(1 + \exp(u))$ is monotonic in u .

(b) Fit a logistic regression model (i.e., minimize $\ell(\theta)$), and also a robust logistic regression model (i.e., minimize $\ell^{\text{wc}}(\theta)$), using the data given in `rob_logistic_reg_data.py`. The x_i s are provided as the rows of a $n \times d$ matrix named `x`. The y_i s are provided as the entries of a n vector named `y`. The file also contains a test data set, `x_test`, `y_test`. Give the test error rate (i.e., fraction of test set data points for which $\hat{y} \neq y$) for the logistic regression and robust logistic regression models.

My Response

We will consider a generalization of this problem. Specifically, we'll actually define the worst-case logistic loss as

$$\ell^{\text{wc}}(\theta) = \sum_{i=1}^n \sup_{\|\delta_i\| \leq \epsilon} \log(1 + \exp(-y_i \theta^T (x_i + \delta_i))),$$

where $\|\delta_i\|$ is just *some norm* on \mathbf{R}^d . We will explore the reasoning behind this generalization after reformulating the model.

The robust logistic regression problem is to solve the unconstrained problem

$$\text{minimize} \quad \ell^{\text{wc}}(\theta) = \sum_{i=1}^n \sup_{\|\delta_i\| \leq \epsilon} \log(1 + \exp(-y_i \theta^T (x_i + \delta_i)));$$

however, while this *is* a convex optimization problem, *it is not* DCP compatible. This is because the supremum operator is an analytical expression; *i.e.*, (for our purposes), it cannot be handled by a computer. Furthermore, we need to reformulate the problem so the supremum operator is neither in the objective function nor in any constraint function. To begin reformulating, we use the hint and move the supremum from the objective function to the constraint set. The corresponding problem is,

$$\begin{aligned} &\text{minimize} \quad \sum_{i=1}^n \log(1 + \exp(u_i)) \\ &\text{subject to} \quad \sup \{-y_i \theta^T (x_i + \delta_i) \mid \|\delta_i\| \leq \epsilon\} \leq u_i, \quad i = 1, \dots, n, \end{aligned}$$

where now both $\theta \in \mathbf{R}^d$ and $u \in \mathbf{R}^n$ are optimization variables. To see that this problem is equivalent to the original unconstrained problem, consider holding θ fixed and optimizing solely over u . Because the log-sum-exp function is monotonic in its input and the objective function is separable, we achieve the optimum over u by choosing

$$u_i = \sup \{-y_i \theta^T (x_i + \delta_i) \mid \|\delta_i\| \leq \epsilon\}$$

for $i = 1, \dots, n$. Furthermore, then optimizing over θ yields the original, unconstrained problem.

Now, as previously mentioned, having the supremum operator in the constraint function set does not make the problem DCP compatible. Furthermore, we continue and finish the reformulation effort by first pulling terms out of the supremum:

$$\sup \{ -y_i \theta^T (x_i + \delta_i) \mid \|\delta_i\| \leq \epsilon \} \leq u_i, \quad i = 1, \dots, n,$$

is equivalent to

$$-y_i \theta^T x_i + \sup \{ -y_i \theta^T \delta_i \mid \|\delta_i\| \leq \epsilon \} \leq u_i \quad i = 1, \dots, n,$$

and then scaling the expression within the supremum to recognize that

$$\sup \{ -y_i \theta^T \delta_i \mid \|\delta_i\| \leq \epsilon \} = \sup \{ -\epsilon y_i \theta^T \delta_i \mid \|\delta_i\| \leq 1 \} = \|\epsilon y_i \theta\|_*,$$

where $\|\cdot\|_*$ is the *dual norm* of whatever norm we choose to define our uncertainty set. Therefore, we've found that our original unconstrained, non-DCP compatible robust logistic regression optimization problem

$$\text{minimize} \quad \ell^{\text{wc}}(\theta) = \sum_{i=1}^n \sup_{\|\delta_i\| \leq \epsilon} \log(1 + \exp(-y_i \theta^T (x_i + \delta_i)))$$

is equivalent to the following constrained, DCP compatible robust logistic regression optimization problem

$$\begin{aligned} &\text{minimize} \quad \sum_{i=1}^n \log(1 + \exp(u_i)) \\ &\text{subject to} \quad -y_i \theta^T x_i + \|\epsilon y_i \theta\|_* \leq u_i, \quad i = 1, \dots, n. \end{aligned}$$

It turns out that the choice of uncertainty set (induced by our choice in norm) dramatically affects the accuracy of this robust model. Table 2 contains the test error rate of a typical logistic regression model versus robust logistic regression models with uncertainty sets induced by ℓ_1 -, ℓ_2 -, and ℓ_∞ -norms and varying values of ϵ . In most every circumstance, at least one, if not all, of the robust models perform better than the vanilla model (and sometimes significantly so). Table 1 contains the raw classification scores of each model for each ϵ for a test data set containing $n = 60$ points.

	$\epsilon = 0.5$	$\epsilon = 0.75$	$\epsilon = 1$
Non-Robust	36	39	35
ℓ_1 Uncertainty	47	44	35
ℓ_2 Uncertainty	49	47	41
ℓ_∞ Uncertainty	37	46	41

Table 1: Number of Correct Classifications out of $n = 60$

	$\epsilon = 0.5$	$\epsilon = 0.75$	$\epsilon = 1$
Non-Robust	0.6	0.65	0.583
ℓ_1 Uncertainty	0.783	0.733	0.583
ℓ_2 Uncertainty	0.817	0.783	0.683
ℓ_∞ Uncertainty	0.617	0.767	0.683

Table 2: Test Error Rate

2 Design and Control

2.1 Minimum Fuel Control

Problem Overview

Imagine you are piloting a commercial plane from NYC to LA. You could fly there in a fuel efficient manner, or you could keep your passengers on their toes and perform an aileron roll as you fly over Arizona.

This problem is simply concerned with finding the control input which takes the system from its starting point to its desired ending point in the most fuel efficient manner.

Problem

[BV04] **Exercise 4.16.** *Minimum fuel optimal control.* Consider the LTI dynamical system with state $x_t \in \mathbf{R}^n$, $t = 0, \dots, N$, and actuator or input signal $u_t \in \mathbf{R}$, for $t = 0, \dots, N - 1$. The dynamics of the system are governed by the linear recurrence

$$x_{t+1} = Ax_t + bu_t, \quad t = 0, \dots, N - 1,$$

where $A \in \mathbf{R}^{m \times n}$ and $b \in \mathbf{R}^n$ are given. Assume the initial state is $x_0 = 0$. The *minimum fuel optimal control problem* is to choose the inputs u_0, \dots, u_{N-1} so as to minimize the total fuel consumed, which is given by

$$F = \sum_{t=0}^{N-1} f(u_t),$$

subject to the constraint that $x_N = x_{\text{des}}$, where N is the (given) time horizon and $x_{\text{des}} \in \mathbf{R}^n$ is the (given) desired final or target state. The function $f : \mathbf{R} \rightarrow \mathbf{R}$ is the *fuel use map* for the actuator, and gives the amount of fuel used as a function of the actuator signal amplitude. In this problem we use

$$f(a) = \begin{cases} |a| & |a| \leq 1 \\ 2|a| - 1 & |a| > 1. \end{cases}$$

Formulate the minimum fuel control problem as an LP.

Response

We want to write this minimum fuel optimal control problem as a LP. However, a LP formulation is rather restrictive, so let's begin by formulating this problem as a convex optimization problem. It is tempting to naively claim that the problem

$$\begin{aligned} & \text{minimize} && \sum_{t=0}^{N-1} f(u_t) \\ & \text{subject to} && x_{t+1} = Ax_t + bu_t, \quad t = 0, \dots, N-1 \\ & && x_0 = 0, \quad x_N = x_{\text{des}}. \end{aligned}$$

is convex. However, we must remember that f is a piecewise function that we are unfamiliar with (*e.g.* it isn't a piecewise linear function). Furthermore, this proposed formulation is not a convex optimization problem. Consider the graph of the fuel use map in figure 4. As stated in the legend of this figure, and clearly seen when observing the graph, the piecewise function f is equivalent to $\max\{|u|, 2|u| - 1\}$. (Of course one could also provide an algebraic argument if so desired.) Furthermore, the optimal fuel control problem can be formulated as

$$\begin{aligned} & \text{minimize} && \sum_{t=0}^{N-1} \max\{|u_t|, 2|u_t| - 1\} \\ & \text{subject to} && x_{t+1} = Ax_t + bu_t, \quad t = 0, \dots, N-1 \\ & && x_0 = 0, \quad x_N = x_{\text{des}}, \end{aligned}$$

which is a valid convex optimization problem. If our goal was simply to find the optimal solution to this fuel problem and we didn't care about formulating the problem as a LP, **we could end our reformulating here.** In fact, the optimal solution $u^* \in \mathbf{R}^{Nm}$, where $Nm = N1 = N$, plotted in figure 5 was obtained by solving this convex problem. Nonetheless, we trudge forward with our LP formulation.

For the time being, to be more concise, let's drop the linear dynamical system constraints

$$x_0 = 0, \quad x_N = x_{\text{des}}, \quad \text{and} \quad x_{t+1} = Ax_t + bu_t, \quad t = 0, \dots, N-1$$

and just consider the unconstrained problem

$$\text{minimize} \quad \sum_{t=0}^{N-1} \max\{|u_t|, 2|u_t| - 1\}.$$

Our first reformulation uses that an objective function defined as the sum of absolute value/maximum expressions can be rewritten as a sum of auxiliary variables, $s \in \mathbf{R}^N$ here, with each summand being less than or equal to an element in the auxiliary variable vector. To simplify the problem further we can also remove the max operator from each summand using that if the maximum element in the set being operated on by max is less than or equal to s_t , then so must every other element. These two reformulation techniques yield the problem

$$\begin{aligned} & \text{minimize} && \mathbf{1}^T s \\ & \text{subject to} && |u_t| \leq s_t, \quad t = 1, \dots, N \\ & && 2|u_t| - 1 \leq s_t, \quad t = 1, \dots, N. \end{aligned}$$




Figure 4: Actuator Fuel Use Map.

To fully linearize the problem, consider $y \in \mathbf{R}^N$ and the problem

$$\begin{aligned}
& \text{minimize} && \mathbf{1}^T s \\
& \text{subject to} && y \preceq s \\
& && 2y - \mathbf{1} \preceq s \\
& && -y \preceq u \preceq y,
\end{aligned}$$

which uses this new variable to “pull out” the absolute value from the two other sets of constraints. Now, before returning the LDS constraints to the formulation observe the pattern in the LDS defining recurrence,

$$\begin{aligned}
x_1 &= Ax_0 + bu_0 \\
x_2 &= Ax_1 + bu_1 \\
&= A(Ax_0 + bu_0) + bu_1 = A^2x_0 + Abu_0 + bu_1 \\
x_3 &= A^3x_0 + A^2b_0 + Abu_1 + bu_2. \\
&\vdots
\end{aligned}$$

Plugging x_0 and x_{des} into the recurrence, the system evolution constraints are introduced back into the optimization problem

$$\begin{aligned}
&\text{minimize} && \mathbf{1}^T s \\
&\text{subject to} && Hu = x_{\text{des}} \\
&&& y \preceq s \\
&&& 2y - \mathbf{1} \preceq s \\
&&& -y \preceq u \preceq y,
\end{aligned}$$

with the equality constraints $Hu = x_{\text{des}}$, where

$$H = [A^{N-1}b \quad A^{N-2}b \quad \dots \quad Ab \quad b].$$

This LP is equivalent to the original minimum fuel optimal control problem, with the corresponding optimal actuator signal plotted in figure 5.




Figure 5: Minimum fuel actuator signal.

2.2 Path Planning with Contingencies

Problem Overview

The problem is to develop three routes for a vehicle. Two routes are “contingency routes,” which address the circumstance that some obstacle impedes the vehicle’s optimal path once the vehicle has already started moving.

Problem

[BV24] **Exercise 16.9.** *Path Planning with contingencies.* A vehicle path down a (straight, for simplicity) road is specified by a vector $p \in \mathbf{R}^N$, where p_i gives the position perpendicular to the centerline at the point ih meters down the road, where $h > 0$ is a given discretization size. (Throughout this problem, indexes on N -vectors will correspond to positions on the road.) We normalize p so $-1 \leq p_i \leq 1$ gives the road boundaries. (We are modeling the vehicle as a point, by adjusting for its width.) You are given the initial two positions $p_1 = a$ and $p_2 = b$ (which give the initial road position and angle), as well as the final two positions $p_{N-1} = c$ and $p_N = d$. You know there may be an obstruction at position $i = O$. This will require the path to either go around the obstruction on the left, which requires $p_O \geq 0.5$, or on the right, which requires $p_O \leq -0.5$, or possibly the obstruction will clear, and the obstruction does not place any additional constraint on the path. These are the three contingencies in the problem title, which we label as $k = 1, 2, 3$. You will plan three paths for these contingencies, $p^{(i)} \in \mathbf{R}^N$ for $i = 1, 2, 3$. They must each satisfy the given initial and final two road positions and the constraint of staying within the road boundaries. Paths $p^{(1)}$ and $p^{(2)}$ must satisfy the (different) obstacle avoidance constraints given above. Path $p^{(3)}$ does not need to satisfy an avoidance constraint. Now we add a twist: You will not learn which of the three contingencies will occur until the vehicle arrives at position $i = S$, when the sensors will determine which contingency holds. We model this with the information constraints (also called causality constraints or non-anticipatory constraints),

$$p_i^{(1)} = p_i^{(2)} = p_i^{(3)}, \quad i = 1, \dots, S,$$

which state that before you know which contingency holds, the three paths must be the same. The objective to be minimized is

$$\sum_{k=1}^3 \sum_{i=2}^{N-1} \left(p_{i-1}^{(k)} - 2p_i^{(k)} + p_{i+1}^{(k)} \right)^2,$$

the sum of the squares of the second differences, which gives smooth paths.

(a) Explain how to solve this problem using convex optimization.

(b) Solve the problem with data given in `path_plan_contingencies_data.*`. The data files include code to plot the results, which you should use to plot (on one plot) the optimal paths. Report the optimal objective value. Give a very brief informal explanation for what you see happening for $i = 1, \dots, S$. Hint. In Python, use the (default) solver ECOS to avoid warnings about inaccurate solutions.

My Response

Fortunately, the problem as stated is already a DCP compatible optimization problem. The corresponding formulation is

$$\begin{aligned}
 & \text{minimize} && \sum_{k=1}^3 \sum_{i=2}^{N-1} \left(p_{i-1}^{(k)} - 2p_i^{(k)} + p_{i+1}^{(k)} \right)^2 \\
 & \text{subject to} && p_i^{(1)} = p_i^{(2)} = p_i^{(3)}, && i = 1, \dots, S \\
 & && p_O^{(1)} \geq 1/2, \\
 & && p_O^{(2)} \leq -1/2, \\
 & && p_1^{(i)} = a, \quad p_2^{(i)} = b, && i = 1, 2, 3 \\
 & && p_{N-1}^{(i)} = c, \quad p_N^{(i)} = d, && i = 1, 2, 3.
 \end{aligned}$$

(b) The code in algorithm 5 is all that is required to solve the problem, and when instatiating

Algorithm 5: Path Planning Code

```

1      p = cp.Variable((3, N))
2
3      constra = [constr for constr_tuple in
4      range(2, S+1)]
5          for constr in constr_tuple]
6      constra += [p[0, 0] >= h,
7                  p[1, 0] <= -h]
8      constra += [p[:, 0] == p1*np.ones(3),
9                  p[:, 1] == p2*np.ones(3),
10                 p[:, N-2] == pN_1*np.ones(3),
11                 p[:, N-1] == pN*np.ones(3)]
12
13     f0 = cp.sum( [ cp.sum_squares( cp.vstack([p[k, i-1] - 2*p[k, i
14 ] + p[k, i+1] for i in range(2, N-2)]) ) for k in range(3)] )
15
16     prob = cp.Problem(cp.Minimize(f0), constra)
17     prob.solve()

```

the model with the provided values, we find $p^* = 0.0001423$ (examining the paths, this low objective value makes sense as there are no sharp deviations along any of the paths). The optimal paths can be found in figure 6. For steps $i = 1, \dots, 60$ we see that indeed all paths follow the same trajectory, but once the object is revealed, the two contingency paths make the necessary adjustments to avoid the object.




Figure 6: Possible Paths.

2.3 Output Tracking

Problem Overview

This problem keeps the dynamics from Problem 2.1 (minimum fuel control), but adds the requirement that the output produced from the system state (so perhaps a reading on the plane's dashboard of its coordinates) matches a predetermined path. In other words, we wish to provide control input to the system in order to track some desired trajectory. There's also an additional twist to the problem solution which would allow this same approach to work in a setting where there's randomness or an infinite time horizon, but these are unimportant for a cursory understanding.

Problem

[Boy-8] **HW7 Q1.MPC for output tracking.** Consider the linear dynamical system

$$x_{t+1} = Ax_t + Bu_t, \quad y_t = Cx_t, \quad t = 0, \dots, T-1,$$

with state $x_t \in \mathbf{R}^n$, input $u_t \in \mathbf{R}^m$, and output $y_t \in \mathbf{R}^p$. The matrices A and B are known, and $x_0 = 0$. The goal is to choose the input sequence u_1, \dots, u_t to minimize the output tracking cost

$$J_{\text{output}} = \sum_{t=1}^T \|y_t - y_t^{\text{des}}\|_2^2,$$

subject to $\|u_t\|_\infty \leq U^{\max}$, $t = 0, \dots, T-1$.

For the remainder of this problem we work with the specific problem instance with associated data

$$A = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}, \quad B = \begin{bmatrix} 0 \\ 0.5 \\ 1 \end{bmatrix}, \quad C = \begin{bmatrix} -1 & 0 & 1 \end{bmatrix},$$

$T = 100$, and $U^{\max} = 0.1$. The desired output trajectory is given by

$$y_t^{\text{des}} = \begin{cases} 0 & t < 30, \\ 10 & 30 \leq t < 70 \\ 0 & t \geq 70. \end{cases}$$

(a) Find the optimal input u^* and the associated optimal cost J^* .

(b) *Rolling look-ahead.* Now consider the input obtained using an MPC-like method where at time t , we find the values of u_t, \dots, u_{t+N-1} that solve the following convex optimization problem

$$\begin{aligned} & \text{minimize} && J_{\text{output}} = \sum_{\tau=t+1}^{t+N} \|Cx_\tau - y_\tau^{\text{des}}\|_2^2 \\ & \text{subject to} && \|u_\tau\|_\infty \leq 0.1, \quad x_{\tau+1} = Ax_\tau + Bu_\tau, \quad \tau = t, \dots, t+N-1 \\ & && x_0 = 0. \end{aligned}$$

The value N is the amount of *look-ahead*, since it dictates how much of the future of the desired output signal we are allowed to access when we decide on the current input.

Find the input signal for look-ahead values $N = 8$, $N = 10$, and $N = 12$. Compare the cost J_{output} obtained in these three instances to the optimal cost J_{output}^* found in part (a).

My Response

(a) This is simply a *linear* (in the dynamics) *time-invariant quadratic tracking* problem. Instead of doing a theoretical analysis of controllability, etc., we can determine the feasibility of the control problem by formulating and attempting to solve the following convex optimization problem:

$$\begin{aligned} & \text{minimize} && J_{\text{output}} = \sum_{t=1}^{100} \|Cx_t - y_t^{\text{des}}\|_2^2 \\ & \text{subject to} && \|u_t\|_{\infty} \leq 0.1, \quad x_{t+1} = Ax_t + Bu_t, \quad t = 0, \dots, 99 \\ & && x_0 = 0, \end{aligned}$$

(with the provided data for A , B , C , and y^{des} , of course.) Using CVXPY, we obtain the optimal cost $J_{\text{output}}^* = 112.4157$.

(b) To solve this problem using a rolling look-ahead, there are two key components. One, we need to solve the same optimization problem as in part (a), but just over a smaller time horizon. Two, we need solve this receding horizon optimization problem at each discretized instance in time, step forward in time corresponding to the first control input in the sequence of N control inputs obtained by solving the optimization problem, and then repeat this loop until we have stepped T steps forward. The corresponding optimization problem can be seen in algorithm 6, and the simulation code can be found in algorithm 7.

The cost obtained by running MPC with

- $N = 8$ is 379.634,
- $N = 10$ is 128.13,
- and $N = 12$ is 123.62.

Figure 7 shows the output trajectories for $N = 8$, $N = 10$, $N = 12$, the optimal output trajectory, and the desired output trajectory.

Algorithm 6: Optimization Problem

```
1         def solve(T, t, x_init):
2             x = cp.Variable((n, T+1))
3             u = cp.Variable((m, T))
4
5             J_output = cp.sum_squares(cp.vstack( [ C@x[:, tau] -
y_des[tau + t - 1, :]
6
7             for tau in range
(1, T+1) ] ))
8
9             obj = cp.Minimize(J_output)
10
11             constr = [x[:, 0] == x_init]
12             for tau in range(T):
13                 constr += [ x[:, tau+1] == A@x[:, tau] + B@u[:,
tau],
14
15                             cp.norm(u[:, tau], 'inf') <= U_max ]
16
17             prob = cp.Problem(obj, constr)
18             prob.solve()
19             return (prob.value, u.value, x.value)
```

Algorithm 7: Model Predictive Control Simulator

```
1         def run_mpc(N, T):
2             x = np.zeros((n))
3             U_traj = np.zeros((m, T))
4             Y = np.zeros((p, T))
5             J_output = 0
6
7             for t in range(T):
8                 _, u, _ = solve(N, t, x)
9                 U_traj[:, t] = u[:, 0]
10                x = A@x + B@u[:, 0]
11                y = C@x
12                Y[:, t] = y
13                J_output += cp.sum_squares(y - y_des[t, :]).value
14
15             return (J_output, Y)
```




Figure 7: Output Trajectories.

2.4 UAV Design

Problem Overview

This is an alternative optimal control problem. More specifically, it is an optimal design problem. While the actual problem prompt is long and contains a lot of moving pieces (which is to be expected given that we wish to design a UAV), it is simple to understand. We want to design an unmanned aerial vehicle. It needs to complete some number of missions where for each mission the distance the UAV must cover is known, the velocity it needs to fly at is known, and the payload that it is carrying is known. The below problem is concerned with designing this UAV such that one, it actually meets the requirements of physics, two, it meets the parameters of each mission (both these prior two are required...they are constraints), all the while minimizing the cost to manufacture and operate the UAV.

Problem

[BV24] **Exercise 18.14.** *Design of an unmanned aerial vehicle.* You are tasked with developing the high-level design for an electric unmanned aerial vehicle (UAV). The goal is to design the least expensive UAV that is able to complete K missions, labeled $k = 1, \dots, K$. Mission k involves transporting a payload of weight $W_k^{\text{pay}} > 0$ (in kilograms) over a distance $D_k > 0$ (in meters), at a speed $V_k > 0$ (in meters per second). These mission quantities are given. The high-level design consists of choosing the engine weight W^{eng} (in kilograms), the battery weight W^{bat} (in kilograms), and the wing area S (in m^2), within the given limits

$$W_{\min}^{\text{eng}} \leq W^{\text{eng}} \leq W_{\max}^{\text{eng}}, \quad W_{\min}^{\text{bat}} \leq W^{\text{bat}} \leq W_{\max}^{\text{bat}}, \quad S_{\min} \leq S \leq S_{\max}.$$

(The lower limits are all positive.) We refer to the variables W^{eng} , W^{bat} , and S as the design variables. In addition to choosing the design variables, you must choose the power $P_k > 0$ (in watts) that flows from the battery to the engine, and the angle of attack $\alpha_k > 0$ (in degrees) of the UAV during mission k , for $k = 1, \dots, K$. These must satisfy

$$0 \leq P_k \leq P_{\max}, \quad 0 \leq \alpha_k \leq \alpha_{\max},$$

where α_{\max} is given, and P_{\max} depends on the engine weight as described below. We refer to these $2K$ variables as the mission variables. The engine weight, battery weight, and wing area are the same for all k missions; the power and angle of attack can change with the mission. The weight of the wing is W^{wing} (in kilograms) is given by $W^{\text{wing}} = C_W S^{1.2}$, where $C_W > 0$ is given. The total weight of the UAV during mission k , denoted W_k , is the sum of the battery weight, engine weight, wing weight, the payload weight, and a baseline weight W^{base} , which is given. The total weight depends on the mission, via the payload weight, and so is subscripted by k . The lift and drag forces acting on the UAV in mission k are

$$F_k^{\text{lift}} = \frac{1}{2} \rho V_k^2 C_L(\alpha_k) S, \quad F_k^{\text{drag}} = \frac{1}{2} \rho V_k^2 C_D(\alpha_k) S$$

(in newtons), where C_L and C_D are the lift and drag coefficients as functions of the angle of attack α_k , and $\rho > 0$ is the (known) air density (in kilograms per cubic meter). We will

use the simple functions

$$C_L(\alpha) = c_L\alpha, \quad C_D(\alpha) = c_{D1} + c_{D0}\alpha^2,$$

where $c_L > 0, c_{D0} > 0$, and $c_{D1} > 0$ are given constants. To maintain steady level flight, the lift must equal the weight, and the drag must equal the thrust from the propeller, denoted T_k (in newtons), i.e.,

$$F_k^{\text{lift}} = W_k, \quad F_k^{\text{drag}} = T_k.$$

The thrust force, power P_k (in watts), and the UAV speed are related via $P_k = T_k V_k$. The engine maximum power is related to its weight by $W^{\text{eng}} = C_P P_{\text{max}}^{0.803}$ where $C_P > 0$ is given. The battery capacity E (in joules) is equal to $C_E W^{\text{bat}}$, where $C_E > 0$ is given. The total energy expended over mission k , with speed V_k , power output P_k , and distance D_k is $P_k D_k / V_k$. This must not exceed the battery capacity E . The overall cost of the UAV is the sum of a design cost and a mission cost. The design cost C_{des} , which is an approximation of the cost of building the UAV, is given by

$$C_{\text{des}} = 100W^{\text{eng}} + 45W^{\text{bat}} + 2W^{\text{wing}}.$$

The mission cost C_{mis} is given by

$$C_{\text{mis}} = \sum_{k=1}^K (T_k + 10\alpha_k),$$

which captures our desire that the thrust and angle of attack be small. In summary, $W_{\text{min}}^{\text{eng}}, W_{\text{max}}^{\text{eng}}, W_{\text{min}}^{\text{bat}}, W_{\text{max}}^{\text{bat}}, S_{\text{min}}, S_{\text{max}}, \alpha_{\text{max}}, W_{\text{base}}, C_W, c_L, c_{D0}, c_{D1}, C_P, C_E$, and ρ are given. Additionally, D_k, V_k , and W_k^{pay} are given for $k = 1, \dots, K$.

(a) The problem as stated is almost a geometric problem (GP). By relaxing two constraints it becomes a GP, and therefore readily solved. Identify these constraints and give the relaxed versions. Briefly explain why the relaxed constraints will be tight at the solution, which means by solving the GP, you've actually solved the original problem. You do not need to reduce the relaxed problem to a standard form GP, or the equivalent convex problem; it's enough to express it in DGP compatible form.

(b) Solve the relaxed problem you formulate in part (a) with data given in the provided python file. Give the optimal costs C_{des}^* and C_{mis}^* , and the values of all design and mission variables. Check that at your solution the relaxed constraints are tight.

My Response

Recall that an optimization problem of the form

$$\begin{aligned} & \text{minimize} && f_0(x) \\ & \text{subject to} && f_i(x) \leq 1, \quad i = 1, \dots, m \\ & && h_i(x) = 1, \quad i = 1, \dots, p, \end{aligned}$$

where f_0, \dots, f_m are posynomials and h_1, \dots, h_p are monomials is a *geometric program* (GP) in *standard form*. (Note that $\mathcal{D} = \mathbf{R}_{++}^n$; the constraint $x \succ 0$ is implicit.)

Most importantly, note that the equality constraint functions must be monomials. This requirement helps direct our search for invalid constraints, *i.e.*, we should look for constraints, which as formulated “verbally,” would require posynomial equalities. The first such constraint(s) is

$$F_k^{\text{lift}} = W_k, \quad k = 1, \dots, K,$$

since

$$\begin{aligned} W_k &= W^{\text{bat}} + W^{\text{eng}} + W^{\text{wing}} + W_k^{\text{pay}} + W^{\text{base}} \\ &= W^{\text{bat}} + W^{\text{eng}} + C_W S^{1.2} + W_k^{\text{pay}} + W^{\text{base}}, \end{aligned}$$

is a *posynomial* and F_k^{lift} is a *monomial*. Because posynomials are *closed under division*, the proposed constraint

$$F_k^{\text{lift}} = W_k \iff (F_k^{\text{lift}})^{-1} W_k = 1, \quad k = 1, \dots, K,$$

is a posynomial equality constraint, which to emphasize again, is invalid for GP formulation.

For the same reason, the constraint(s)

$$F_k^{\text{drag}} = T_k, \quad k = 1, \dots, K,$$

is also invalid. However, in this case, it is the force term, F_k^{drag} , which is the posynomial and $T_k = (1/V_k)P_k$ which is the monomial. To formulate a proper GP, we therefore make the following two (or more accurately, $2K$) relaxations

$$W_k = F_k^{\text{lift}}, \quad \text{and} \quad F_k^{\text{drag}} = T_k, \quad k = 1, \dots, K$$

become

$$W_k \leq F_k^{\text{lift}}, \quad \text{and} \quad F_k^{\text{drag}} \leq T_k, \quad k = 1, \dots, K,$$

or equivalently (and in GP standard form)

$$(F_k^{\text{lift}})^{-1} W_k \leq 1, \quad \text{and} \quad T_k^{-1} F_k^{\text{drag}} \leq 1, \quad k = 1, \dots, K,$$

in the relaxed formulation.

It is now highly tempting to start writing down an optimization problem by just reading off the problem specifications. *This is discouraged.* Even though we have found and proposed relaxations for the two originally invalid GP constraints, it is still *very easy* to formulate a *non-DGP compatible problem*. As an example, a common (and understandable) mistake would be to include the constraints

$$0 \leq P_k, \quad k = 1, \dots, K,$$

in the formulation. However, **including them** would make the formulation **non-DGP compatible**. All constraints, equality or inequality, must have a 1 on the right-hand side of the constraint when reduced to standard form. The constraints $0 \leq P_k$, $k = 1, \dots, K$ cannot be reduced to this form. However, these are valid GP constraints! They are just

included *implicitly* since the domain of the problem is \mathbf{R}_{++}^n . Including them *explicitly* leads to a non-DGP compatible formulation.

Additionally, unlike in LPs, where additional “linking” constraints are harmless, adding such syntactic constraints here can, again, lead to a non-DGP compatible problem. As an example, consider the specification that the lift force acting on the UAV in mission k is

$$F_k^{\text{drag}} = \frac{1}{2} \rho V_k^2 C_D (\alpha_k) S,$$

which as previously discussed is a *posynomial*. In an LP formulation (at least, as commonly taught in an operations research setting), it is encouraged (and would be totally valid) to declare F_k^{drag} as an optimization variable, include the above equality in the LP **as a constraint**, and then use F_k^{drag} throughout the remainder of the problem in constraints involving drag force. However, this “linking” constraint is **non-DGP compatible** because it would be reduced to a posynomial equality constraint, which as discussed at great length above, is **not a valid GP constraint**.

Furthermore, we must distinguish between actual *optimization variables* and what I’ll call *derived optimization variables*. The former are what they are always defined as: the variables that are to be chosen to minimize the objective function. We will define the latter as a mathematical expression whose value is decided by the optimization variables. To (ab)use a Computer Science dialect, the derived optimization variables are functions of references to the optimization variables (where the function can take multiple optimization variables and even other derived optimization variables). Consequently, we can use either the derived variables or the actual variables when formulating our problem. As an example, the constraint

$$\frac{1}{2} \rho V_k^2 C_D (\alpha_k) S \leq (1/V_k) P_k, \quad k = 1, \dots, K$$

is equivalent to

$$F_k^{\text{drag}} \leq T_k, \quad k = 1, \dots, K.$$

However, $F_k^{\text{drag}} = \frac{1}{2} \rho V_k^2 C_D (\alpha_k) S$ and $T_k = (1/V_k) P_k$, $k = 1, \dots, K$, should **not** be included in the problem constraints and F_k^{drag} , T_k , $k = 1, \dots, K$ **are not** optimization variables. They both are merely expressions *encoding* optimization variables. As a final attempt to stress the difference, when building this problem with a declarative language such as CVXPY in Python, the problem should be created as follows (and see Algorithm 8 for the actual code)

1. Declare optimization variables. Ex. `opt_var = cvxpy.Variable(num_entries)` (notice the explicit use of the `Variable` keyword).
2. Create derived optimization variables. Ex. `der_var = 2.5*opt_var` (this becomes a `cvxpy.Expression`).
3. Create the constraints and objective using the optimization variables, the derived optimization variables, or some combination of the two. Ex. `der_var <= 5`, which would be equivalent to adding the constraint `opt_var <= 2`.
4. Solve the problem.

This is a nuanced point, and perhaps is obvious to readers who do not have training in OR formulating LPs, but it is critically important.

We now put everything together; the following is the Relaxed UAV Design Problem.

Optimization Variables

- Design Variables: $W^{\text{eng}}, W^{\text{bat}}, S$.
- Mission Variables: $P_k, \alpha_k, \quad k = 1, \dots, K$.

Derived Variables/CVXPY expressions

- $W^{\text{wing}} = C_W S^{1.2}$,
- $W_k = W^{\text{bat}} + W^{\text{eng}} + W^{\text{wing}} + W_k^{\text{pay}} + W^{\text{base}}$,
- $F_k^{\text{lift}} = \frac{1}{2} \rho V_k^2 C_L(\alpha_k) S, \quad k = 1, \dots, K$,
- $F_k^{\text{drag}} = \frac{1}{2} \rho V_k^2 C_D(\alpha_k) S, \quad k = 1, \dots, K$,
- $P_{\max} = (C_P^{-1} W^{\text{eng}})^{1/0.803}$,
- $T_k = (1/V_k) P_k, \quad k = 1, \dots, K$,
- $E = C_E W^{\text{bat}}$,
- $C_{\text{des}} = 100 W^{\text{eng}} + 45 W^{\text{bat}} + 2 W^{\text{wing}}$,
- $C_{\text{des}} = \sum_{k=1}^K (T_k + 10 \alpha_k)$.

Relaxed Formulation

$$\begin{aligned}
& \text{minimize} && C_{\text{des}} + C_{\text{mis}} \\
& \text{subject to} && W_{\min}^{\text{eng}} \leq W^{\text{eng}} \leq W_{\max}^{\text{eng}}, \quad W_{\min}^{\text{bat}} \leq W^{\text{bat}} \leq W_{\max}^{\text{bat}}, \quad S_{\min} \leq S \leq S_{\max} \\
& && P_k \leq P_{\max}, \quad \alpha_k \leq \alpha_{\max}, \quad k = 1, \dots, K \\
& && F_k^{\text{lift}} \geq W_k, \quad k = 1, \dots, K \\
& && F_k^{\text{drag}} \leq T_k, \quad k = 1, \dots, K \\
& && (1/V_k) P_k D_k \leq E, \quad k = 1, \dots, K.
\end{aligned}$$

Problem Solution and Analysis.

Firstly, we address why the relaxed constraints, $F_k^{\text{lift}} \geq W_k$ and $F_k^{\text{drag}} \leq T_k$, $k = 1, \dots, K$, will be tight at the solution.

The tightness in the constraint $F_k^{\text{drag}} \leq T_k$ at the solution is easy to see. We proceed with a semi-formal argument. Suppose that the drag inequality for mission k holds strictly. That is, $F_k^{\text{drag}} < T_k$. If this is the case, then we can reduce the thrust for mission k by decreasing the power, P_k , that flows from the battery to the engine during that mission as T_k strictly increases and strictly decreases as P_k increases or decreases. We are able to decrease the power because the only lower bound on P_k is the implicit lower bound $P_k \geq 0$, and if it is the case that $P_k \rightarrow 0$, then F_k^{drag} must also go to zero since it too is positive. As both terms

go to zero, we end up violating the original assumption that F_k^{drag} is strictly less than the thrust. Furthermore, being able to reduce T_k contradicts the assumption that we are at the solution since shrinking T_k lowers the mission cost, and thus the overall cost.

A nearly identical argument can be made for the tightness in the constraint F_k^{lift} with respect to the attack angle optimization variable α_k .

(b) The Python code in Algorithm 8 contains the core Python code required to compute the optimal UAV design (it doesn't include constant instantiations and helper functions). The corresponding optimal values are

- $C_{\text{des}}^* = 4449.324$,
- $C_{\text{mis}}^* = 1556.587$,
- $W^{\text{eng}} = 9.951 \text{ kg}$,
- $W^{\text{bat}} = 74.06 \text{ kg}$,
- $S = 7.999 \text{ m}^2$,
- the power values are

$$\begin{aligned} P_1 &= 40.159 \text{ kW}, & P_2 &= 12.142 \text{ kW}, & P_3 &= 20.018 \text{ kW}, \\ P_4 &= 24.998 \text{ kW}, & P_5 &= 8.080 \text{ kW}, \end{aligned}$$

- and the attack angles are

$$\alpha_1 = 0.1377^\circ, \quad \alpha_2 = 0.3141^\circ, \quad \alpha_3 = 0.2249^\circ, \quad \alpha_4 = 0.2114^\circ, \quad \alpha_5 = 0.3785^\circ.$$

Finally, in table 3 we see that the relaxed constraints are indeed tight.

k	T_k	F_k^{drag}	F_k^{lift}	W_k
1	489.739 N	489.739 N	269.648 N	269.646 N
2	220.763 N	220.763 N	276.647 N	276.646 N
3	307.969 N	307.969 N	276.648 N	276.646 N
4	357.119 N	357.119 N	301.648 N	301.646 N
5	168.332 N	168.332 N	253.946 N	253.946 N

Table 3: Tightness in Relaxed Constraints.

Algorithm 8: UAV Design

```
1      ### optimization variables ###
2      W_eng = cp.Variable(1, pos=True)
3      W_bat = cp.Variable(1, pos=True)
4      S = cp.Variable(1, pos=True)
5      P = cp.Variable(K, pos=True)
6      alpha = cp.Variable(K, pos=True)
7
8      ### derived expressions ###
9      W_wing = CW*S**(1.2)
10     W_k = [W_bat + W_eng + W_wing + W_base + W_pay[k] for k in
range(K)]
11     P_max = ((1/CP)**(1/0.803)) * (W_eng**(1/0.803))
12     F_k_lift = [0.5*rho* V[k]**2 * C_L(alpha[k]) * S for k in
range(K)]
13     F_k_drag = [0.5*rho * V[k]**2 * C_D(alpha[k])*S for k in range
(K)]
14     T_k = [P[k] * (1/V[k]) for k in range(K)]
15     E = CE*W_bat
16     C_des = 100*W_eng + 45*W_bat + 2 * W_wing
17     C_mis = cp.sum([T_k[k] + 10*alpha[k] for k in range(K)])
18     ### ###
19
20     ### constraints and objective ###
21     constraints = [W_eng_min <= W_eng, W_eng <= W_eng_max,
22                   W_bat_min <= W_bat, W_bat <= W_bat_max,
23                   S_min <= S, S <= S_max]
24     constraints += [constr for constr_tuple in [(P[k] <= P_max,
25         alpha[k] <= alpha_max) for k in range(K)]
26                   for constr in constr_tuple]
27     constraints += [F_k_lift[k] >= W_k[k] for k in range(K)]
28     constraints += [F_k_drag[k] <= T_k[k] for k in range(K)]
29     constraints += [P[k] * D[k] * (1/V[k]) <= E for k in range(K)]
30
31     obj = cp.Minimize(C_des + C_mis)
32     prob = cp.Problem(obj, constraints)
33     # print(prob.is_dgp())
34     prob.solve(gp=True)
35
```

2.5 Optimal Spacecraft Landing

Problem Overview

I recommend watching the videos linked below. This is essentially a “toy” version of what SpaceX does. (Although, it is worth noting that these solution techniques are “in the family” of the techniques SpaceX uses to land rockets; See [\[All21\]](#).)

Problem

[\[BV24\]](#) **Exercise 16.2.** *Optimal Spacecraft Landing.* We consider the problem of optimizing the thrust profile for a spacecraft to carry out a landing at a target position. The spacecraft dynamics are

$$m\ddot{p} = f - mge_3,$$

where $m > 0$ is the spacecraft mass, $p(t) \in \mathbf{R}^3$ is the spacecraft position, with 0 the target landing position and $p_3(t)$ representing height, $f(t) \in \mathbf{R}^3$ is the thrust force, and $g > 0$ is the gravitational acceleration. (For simplicity we assume that the spacecraft mass is constant. This is not always a good assumption, since the mass decreases with fuel use. We will also ignore any atmospheric friction.) We must have $p(T^{\text{td}}) = 0$ and $\dot{p}(T^{\text{td}}) = 0$, where T^{td} is the touchdown time. The spacecraft must remain in a region given by

$$p_3(t) \geq \alpha \|(p_1(t), p_2(t))\|_2,$$

where $\alpha > 0$ is a given minimum glide slope. The initial position $p(0)$ and velocity $\dot{p}(0)$ are given.

The thrust force $f(t)$ is obtained from a single rocket engine on the spacecraft, with a given maximum thrust; an attitude control system rotates the spacecraft to achieve any desired direction of thrust. The thrust force is therefore characterized by the constraint $\|f(t)\|_2 \leq F^{\text{max}}$. The fuel use rate is proportional to the thrust force magnitude, so the total fuel use is

$$\int_0^{T^{\text{td}}} \gamma \|f(t)\|_2 dt$$

where $\gamma > 0$ is the fuel consumption coefficient. The thrust force is discretized in time, i.e., it is constant over consecutive time periods of length $h > 0$, with $f(t) = f_k$ for $t \in [(k-1)h, kh)$, for $k = 1, \dots, K$, where $T^{\text{td}} = Kh$. Therefore we have

$$v_{k+1} = v_k + (h/m)f_k - hge_3, \quad p_{k+1} = p_k + (h/2)(v_k + v_{k+1}),$$

where p_k denotes $p((k-1)h)$, and v_k denotes $\dot{p}((k-1)h)$. We will work with this discrete-time model. For simplicity, we will impose the glide slope constraint only at the times $t = 0, h, 2h, \dots, Kh$.

(a) *Minimum fuel descent.* Explain how to find the thrust profile f_1, \dots, f_K that minimizes fuel consumption, given the touchdown time $T^{\text{td}} = Kh$ and discretization time h .

(b) *Minimum time descent.* Explain how to find the thrust profile that minimizes the touchdown time, i.e., K , with h fixed and given. Your method can involve solving several convex optimization problems.

(c) Carry out the methods described in parts (a) and (b) above on the problem instance with data given in `spacecraft_landing_data.py`. Report the optimal total fuel consumption for part (a), and the minimum touchdown time for part (b). The data files also contain plotting code (commented out) to help you visualize your solution. Use the code to plot the spacecraft trajectory and thrust profiles you obtained for parts (a) and (b).

Remarks. If you'd like to see the ideas of this problem in action, watch these videos:

- [Grasshopper Diver, Single Cam](#)
- [Grasshopper 24-story Hover](#)
- [Falcon 9 Landing at Zone 1](#)
- [Falcon 9 First Stage Landing](#)

My Response

(a). Formulating this problem is very straightforward as all control specifications are already DCP compatible. After discretizing the integral expression for total fuel use as

$$J_{\text{input}} = \sum_{k=0}^{K-1} \gamma \|f_k\|_2,$$

the minimum fuel descent problem can immediately be formulated as

$$\begin{aligned} & \text{minimize} && J_{\text{input}} \\ & \text{subject to} && v_{k+1} = v_k(h/m)f_k - hge_3, && k = 0, \dots, K-1 \\ & && p_{k+1} = p_k + (h/2)(v_k + v_{k+1}), && k = 0, \dots, K-1 \\ & && \alpha \left\| \begin{bmatrix} (p_k)_1 \\ (p_k)_2 \end{bmatrix} \right\|_2 \leq (p_k)_3, && k = 1, \dots, K \\ & && \|f_k\|_2 \leq F^{\max}, && k = 0, \dots, K-1 \\ & && v_0 = \dot{p}(0), \quad p_0 = p(0) \\ & && v_K = 0, \quad p_K = 0. \end{aligned}$$

(Note that we've adopted a Pythonic zero-based indexing for the formulation, i.e., the first force we apply is at $t = 0$, right when we start controlling the rocket).

(b). Keeping the parameter h fixed, we reduce the time it takes to descend by decreasing K , the number of steps it takes to move from $(p(0), \dot{p}(0)) \rightarrow (0, 0)$. The simplest way to find

K is with a linear search where we start by solving the feasibility problem,

$$\begin{aligned}
& \text{minimize} && 0 \\
& \text{subject to} && v_{k+1} = v_k(h/m)f_k - hge_3, && k = 0, \dots, K-1 \\
& && p_{k+1} = p_k + (h/2)(v_k + v_{k+1}), && k = 0, \dots, K-1 \\
& && \alpha \left\| \begin{bmatrix} (p_k)_1 \\ (p_k)_2 \end{bmatrix} \right\|_2 \leq (p_k)_3, && k = 1, \dots, K \\
& && \|f_k\|_2 \leq F^{\max}, && k = 0, \dots, K-1 \\
& && v_0 = \dot{p}(0), \quad p_0 = p(0) \\
& && v_K = 0, \quad p_K = 0,
\end{aligned}$$

with a K which yields a feasible solution, then solving the same feasibility problem again but with $K^{\text{curr}} = K - 1$. We continue to decrement the planning horizon, K^{curr} by one until the problem becomes infeasible. At this point, we know that $K^{\text{curr}} + 1$ and the associated solution vector $f^* = (f_1, \dots, f_{K^{\text{curr}}+1})$ is the minimum number of time steps and associated thrust profile which minimizes the touchdown time.

(c). For part (a), the total fuel consumption is about 193.0. For part (b), we find that $K = 25$ minimizes the touchdown time. The minimum fuel descent trajectory and associated thrust profiles are shown in figure 8 and the minimum touchdown time trajectory and associated thrust profiles are shown in figure 9.




Figure 8: Minimum Fuel Descent Trajectory and Thrust Profiles.




Figure 9: Minimum Touchdown Descent Trajectory and Thrust Profiles.

References

- [All21] Stanford SystemX Alliance. *Stephen Boyd's CVXGEN software helps guide SpaceX Falcon*. 2021. URL: <https://systemx.stanford.edu/news/2021-01-05-000000/stephen-boyd%E2%80%99s-cvxgen-software-helps-guide-spacex-falcon>.
- [Boy-8] Stephen Boyd. *EE36b: Convex Optimization II*. 2007-8. URL: <https://see.stanford.edu/Course/EE364B>.
- [BV04] Stephen Boyd and Lieven Vandenberghe. *Convex Optimization*. Cambridge, UK: Cambridge University Press, 2004. ISBN: 978-0521833783.
- [BV24] Stephen Boyd and Lieven Vandenberghe. *Additional Exercises for Convex Optimization*. 2024. URL: https://github.com/cvxgrp/cvxbook_additional_exercises.